

Digit-Level House Number Recognition

Using Convolutional Neural Network

Peter Lai
October 31st, 2016

I. Definition

Overview

Vision-based recognition system is fundamental in building modern AI agent. Information in pixel is far more abundant than any other kinds of carriers. To teach computer to see is one of the coolest challenges nowadays. The most popular method people use in building a visual cognitive system is to train deep [convolutional neural networks](#) on various datasets.

In this project I will use the [Street View House Number](#) or SVHN datasets to build a natural digit recognition model using convolutional neural network. House number recognition is essential in automatic map making and intelligent travel technologies like self driving car. It can also be seen as an introduction to word-level text recognition and interpretations.

The data used in this project are images of cropped house number digits captured in daily street view. Detailed information on the datasets will be discussed in part II.

Problem Statement

This project will focus on the single digit recognition. Given a cropped part of a door plate, the model should be able to recognize what number it is (from 0-9). For example, for a cropped image like **Fig. 1**, the model will return 3 because 3 is center aligned in the image.




Fig. 1

The model can be represented as

$$d = \operatorname{argmax} F(X)$$

Given a image X , model F should return a vector of probability, the prediction digit d is the index where give the highest probability.

The ideal solution is obtained when $d == y$ (the label) is satisfied for every image in the test set and a better solution is the one that makes more correct classification on the test set.

To obtain the solution, the strategy is to minimize the softmax cross entropy on the output of a convolutional neural network using Stochastic Gradient Descent. These can be easily implemented using Google's open source deep learning framework, [TensorFlow](#).

In TensorFlow, nodes have both forward-propagation and back-propagation functionalities. Loss function and iterative optimizer are provided as training agents in a computational graph, which holds all the nodes and operations in a network. All the data is flowing in these nodes according to the rules set by operations and simultaneously optimizing the parameters against the cross entropy.

An architecture worth trying is [LeNet-5](#), the architecture is shown in **Fig. 2**. However, the SVHN data is more complicated, therefore increasing the convolutional layers' depth might be necessary. In addition, the model's output layer should be modified to be a softmax then argmax layer.




Fig. 2

Metrics

1. **Accuracy** (to maximize), is the ratio of how many labels the model correctly recognizes in the dataset.

2. **Precision** and **Recall** (to maximize), will be presented in a confusion matrix. For predictions pointing to a certain class, precision is the ratio of how many images that truly belong to this class. For images belong to a certain class, recall is the ratio of how many images that are successfully recognized. The confusion matrix for the problem looks like **Table 1**. In convenience, I just show the precision and recall for 0 and 3.

Prediction Class	0	1	2	3	4	5	6	7	8	9	Recall
0	200	3	8	3	2	12	1	3	41	3	0.72
1	3	*	*	2	*	*	*	*	*	*	*
2	5	*	*	1	*	*	*	*	*	*	*
3	6	5	2	130	4	4	2	2	1	2	0.82
4	8	*	*	4	*	*	*	*	*	*	*
5	3	*	*	6	*	*	*	*	*	*	*
6	0	*	*	7	*	*	*	*	*	*	*
7	2	*	*	0	*	*	*	*	*	*	*
8	1	*	*	1	*	*	*	*	*	*	*
9	3	*	*	3	*	*	*	*	*	*	*
Precision	0.87	*	*	0.83	*	*	*	*	*	*	

Table 1

3. **Cross Entropy**, also known as Softmax loss (to minimize) the loss function used in training, measuring how close the model is from the ideal model. The cross entropy can be represented as :

$$\text{Cross_Entropy} = (\sum_i \ln(e^{F(x)(y)} / \sum e^{F(x)})) / n$$

where i is the index for images, n is the number of images and F(x)(y) is the probability for the image X to belong to class y.

The cross entropy is not the a metric in evaluate the model. It is the method of calculating the loss in training.

II. Analysis

Data Exploration and Visualization

1. Data Format :

Data was downloaded from <http://ufldl.stanford.edu/housenumbers/>. The website has two categories of data. The dataset used in this project was stored in .mat files, named as "train.mat", "test.mat" and "extra.mat". Two ndarrays will be returned by reading each file using Scipy.io.loadmat() function.

The first ndarray one (name it X) stands for the images, the second (name it y) contains the class labels correspond to X with 0 marked as 10 (MATLAB style).

X is of shape :

(pixel width, pixel height, number of channels, number of items)

y is of shape :

(number of items,)

Table 2 shows the number of images in each file.

File Name	train_32x32.mat	test_32x32.mat	extra_32x32.mat
Number of Record	73257	26032	531131

Table 2

The images in this dataset is in same size 32x32x3, where 3 represent RGB channels.

2. Distribution :

In the Jupiter notebook `model/preprocess_mat.ipynb`, the first few cells will explore the data and plot the distributions of the class labels in all three datasets using function

`peek_distribution` in `model/display_center.py`

See Fig. 3.1, 3.2, 3.3 :

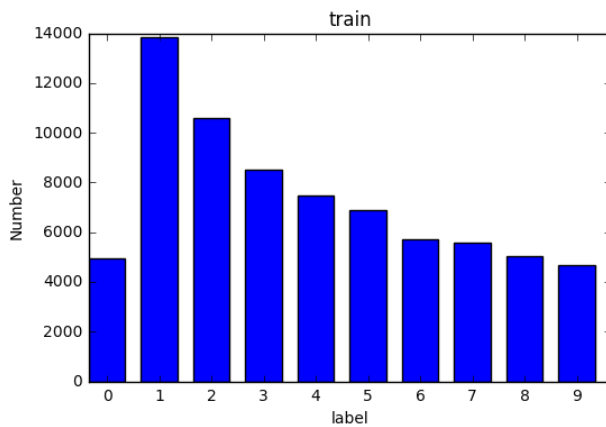


Fig. 3.1




Fig. 3.2




Fig. 3.3

Algorithms and Techniques

Frameworks and Tools

- The core scripts are presented in [Jupyter Notebook](#) using python.
- Scientific Computation & Visualization : [Numpy](#), [Matplotlib](#).
- [scikit-learn](#): A machine learning library that implements classic machine learning algorithms.
- [Tensorflow](#): This is a framework that help users to build computational graph with various kinds of nodes that have functionality of both forward and backward propagating computation, which make it easier to construct deep learning architecture.

Mathematics

- Convolution Neural Network

a. Feature Extractor

Convolution Neural Networks(Fukushima, 1980; LeCun et al., 1988) are neural networks that contain convolution layer. Intuitively, convolution layers are layers that divides a image into small pieces by a windows which moves both horizontally and vertically by a certain stride of pixels. Each piece is seen as a feature, for example, the sharp head of the letter 'A'. Those feature, as group of pixel values, will then be dot product with a group of weight and become a value for the next layer. See **Fig. 4**

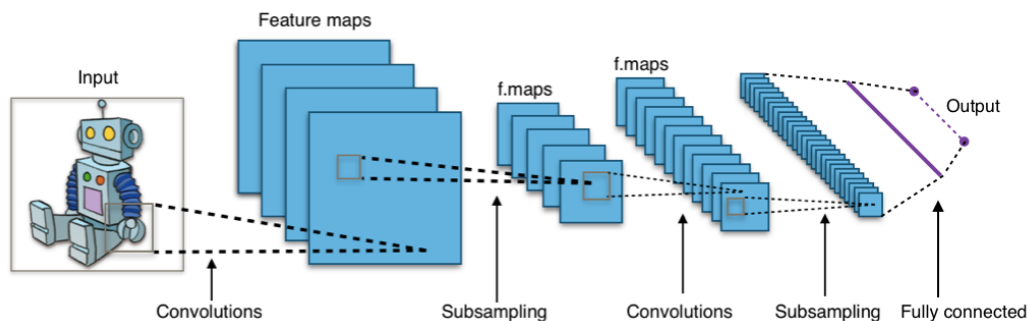


Fig. 4

source: https://en.wikipedia.org/wiki/Convolutional_neural_network#/media/File:Typical_cnn.png
the idea of convolutional neural networks is to extract features

b. Activation Function

Activation is an important step for neural network. It should stop the propagation when a neuron is not activated, which is simply whether the numbers is positive. There are lots of activation functions, like step, sigmoid and tanh, but most of them have various problems. In practice, people use rectified linear unit, ReLU right behind convolution. This unit or its improved version are popular because it's easy to both forward-propagate and back-propagate. See **Fig. 5**, the blue line represents the rectified linear unit.




Fig. 5

[https://en.wikipedia.org/wiki/Rectifier_\(neural_networks\)](https://en.wikipedia.org/wiki/Rectifier_(neural_networks))

c. Pooling and Connecting to Fully Connected Layers

Convolution layers gives output in unfriendly image like shape. Before taking it as the input to create one hot encoding vector, they should be flattened into stripe like vectors. But the output could be very large when images with many depths are flattened. Usually, people apply a method called pooling. Like feature extraction, a small window was moving horizontally and vertically on the image and simply take an average or maximum inside the window. Pooling can help reduce the number of weights to train in the architecture because pooling doesn't need any parameter.

- Training

Training a deep neural network with convolution layers take more techniques than traditional fully connected neural networks. For me there are two basic ideas should be kept in mind: first is to make the information flow efficiently, second is when you are trying to go to the global minimum, adjust your pace according to the terrain.

a. Initialization of Weights

The rule “make it random small” no more applies to deep neural networks. This is very intuitive: if you make all the weights random small, when the information, or numbers are back-propagating, the minus power is accumulating, therefore the numbers could become zero when it arrive some deep neurons and stop the back-propagation, which means the upstream layers will never be trained. In this project, I simply increase the standard deviation of the distribution from upstream to downstream to avoid the issue.

b. Dropout

Dropout is to block the numbers flowing forward and backward in some neuron randomly. It can be seen as an ensemble technique, because in every step, only a part of the network is trained. However, dropout is not implemented in validation and test. It's the equivalence to combine those partial model. And most of the cases, this technique will give more accuracy.

c. Stochastic Gradient Descent

When it comes to large dataset, it would be too heavy for computer to use all the data in one iteration. One solution is to use Stochastic Gradient Descent **** cite **** Basically, we take a mini batch of the dataset and feed it to the network. It may take more iteration for the model to converge, but it works in practice.

d. Update and Optimizer

Basically, the update space is a bowl shape with steep and smooth directions. In this project, I applied Adam update (Kingma and Ba, 2014), which makes the update more on the steep direction.

Adam optimization can be represented as follows: (in python)

```
m = beta1*m + (1-beta1)*dx
v = beta2*v + (1-beta2)*(dx**2)
x += - learning_rate * m / (np.sqrt(v) + eps)
```

source : <http://cs231n.github.io/neural-networks-3/#ada>

Detailed discussion on the mathematics behind the formula would be long, so it's skipped.

- Logging and Visualization

In experiment, I simply use print function to dynamically see the numerical value of the loss and the accuracy.

When the model is satisfactory, I used Tensorflow built-in TensorBoard to visualize the architecture and the tracing the loss and accuracy.

III. Methodology

Data Preprocessing

I took roughly four steps to preprocess the data in `model/preprocess_mat.ipynb`:

1. Reshape the Data

The data is not friendly for Numpy, so I just made the index of items the first dimension. The shape after this step:

```
(number of items, pixel width, pixel height, number of channels)
(number of items, 32, 32, 3)
```

2. Transfer the Images in Grayscale

Since the task is to recognize the digits' profile, grayscale is enough. I applied mean grayscale in this project.

$$\text{grayscale_val} = (R + B + G) / 3$$

3. GCN the Images

Global contrast normalized the images. The approach I apply is the same as what this describes : <http://dp.readthedocs.io/en/latest/preprocess/>

For each image, I subtract the mean of the pixels' gray scale. Then I divide every pixel by the standard deviation across the grayscale values.

After the GCN method, the data become zero-means and the values are scaled down.

Here is an example of a image before and after the GCN transformation:


before :

```
[[[ 33.66666667 19.        17.        ..., 64.33333333 72.        75.66666667]
 [ 34.        20.33333333 17.        ..., 43.        59.        73.66666667]
 [ 39.66666667 21.66666667 18.66666667 ..., 32.33333333 40.33333333 64.66666667]
 ...,
 [ 80.66666667 80.        75.33333333 ..., 98.33333333 96.66666667 95.33333333]
 [ 78.        79.        76.66666667 ..., 99.        98.66666667 98.66666667]
 [ 80.33333333 81.        77.        ..., 100.66666667 100.        96.        ]]
```


after :

```
[[[-1.28437269 -1.91528671 -2.00132044 ..., 0.03481117 0.36460714 0.52233565]
 [-1.27003373 -1.85793089 -2.00132044 ..., -0.88288195 -0.19461211 0.43630191]
 [-1.0262715  -1.80057507 -1.92962567 ..., -1.34172851 -0.99759359 0.04915013]
 ...,
 [ 0.73741997 0.70874206 0.50799669 ..., 1.49738459 1.42568981 1.36833399]
 [ 0.62270833 0.6657252 0.56535251 ..., 1.5260625 1.51172354 1.51172354]
 [ 0.72308102 0.75175893 0.57969147 ..., 1.59775727 1.56907936 1.3970119 ]]
```

before :



after :



4. Collect a Normal Distributed Training Set and Validation Set

As the data exploration shown, real world data has more 1s and 2s than other numbers.

In training, the model should have equal chance to meet all ten labels. Therefore, I collect data from both `train_32x32.mat` and `extra_32x32.mat` to build my own training dataset of normal distribution. I also collect a normal distributed validation set, using the method shown in this paper

5. Shuffle the Data

I simply randomize the index and recollect the data.

Benchmark

A benchmark model I consider is the [**Conv. Maxout + dropout**] architecture used by Goodfellow et al.. Concretely, the model consists of three convolution & maxout layers and a fully connected maxout layer, followed by a fully connected softmax layer. The best accuracy on the test set the team achieved then was 97.53%. (link: <https://arxiv.org/pdf/1302.4389v4.pdf>)

Implementation

(Note: conv means conv followed by relu, fc means fully connected)

Due to hardware limitation, my strategy to gradually scale up my model from a simple **[conv – fc – fc]** network.

In this project I use Stochastic Gradient Descent to do the update. I equally slice the shuffled dataset into mini batches and use them sequentially. The update strategy is Adam update (Kingma and Ba, 2014).

The implementation workflow is shown in `model/cnn_mat.ipynb`. Workflow is briefly three steps :

1. Unpickle the .pickle file created by `model/preprocess_mat.ipynb` and store the training, validation and testing datasets.
2. Build a TensorFlow computational graph and Run in a TensorFlow session.
3. Check the result according to the metric discussed before.

Refinement

Overview of this part

The `model/experiment.ipynb` contains all the milestone model I've tried. Due to hardware limitation, I didn't fine tune some hyper parameters like batch size, bias initialization and training steps. Instead I focus on architecture, weights initialization and update manners with these parameters unchanged.

The title of a experiments illustrate the key factor I focus on in that certain experiment. For example, if an experiment is named 'add another convolutional layer', I mean the experiment has all the other factor tuned to make the deeper network works, but they are less important in my non-demonstrated experiments.

Experiments

Here are some milestone model I've tried :

(Note: conv means conv followed by relu, fc means fully connected, validation set only used for intermediate testing)

1. The Original Model :

architecture :

conv (zero padding)	5 x 5 x 32, stride 2
fc	8192 x 128
fc	128 x 10

weights initialization : TensorFlow default

bias initialization : 0.0 for conv, 1.0 for fc

learning rate : 1e-3

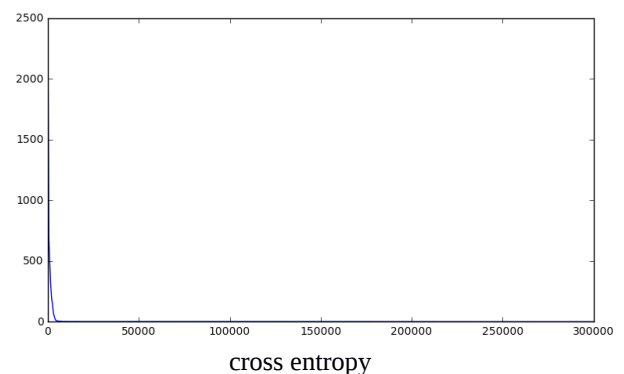

batch size : 256

steps : 30k

result :

a. test accuracy: 71.09%

b. learning curve:



comment : The learning rate is not good for the model. Training curve looks funny.

The cross entropy drops very sharply in the starting phase. The plot won't tell, but actually the cross entropy did not become lower than 1.0 within 30k steps.

c. confusion matrix

Confusion Matrix of testing data:

```
[[1245  28  60  15  63  10 127  2  51 143]
 [ 111 3786 174  55 620  9 18 280 27 19]
 [  36  217 3099 179 103  7 13 340 26 129]
 [  20  58  192 1781 161 151 39 132 183 165]
 [  73 160  37  71 1935 18 136 37 30 26]
 [  16  15  17 361 121 1580 127 13 120 14]
 [  85  20  4  42 171 148 1372 7 114 14]
 [   8  95 116 102 56  5  2 1622 10  3]
 [  71  12  25  78  78 36 149  7 1019 185]
 [ 114  13 174  52  41  3  17  8 105 1068]]
```

Classification report of testing data:

	precision	recall	f1-score	support
0	0.70	0.71	0.71	1744
1	0.86	0.74	0.80	5099
2	0.80	0.75	0.77	4149
3	0.65	0.62	0.63	2882
4	0.58	0.77	0.66	2523
5	0.80	0.66	0.73	2384
6	0.69	0.69	0.69	1977
7	0.66	0.80	0.73	2019
8	0.60	0.61	0.61	1660
9	0.60	0.67	0.64	1595
avg/total	0.72	0.71	0.71	26032

comment: See the conclusion

conclusion :

As shown in the confusion matrix, there are 620 '1's that are classified as '4's. Intuitively, the reason might be that the model cannot recognize the little triangle of '4' and takes it as '1'. It seems the model doesn't extract enough feature now. Therefore more convolutional layers might help.

2. One More Convolutional Layer

architecture :

conv (zero padding)	5 x 5 x 32, stride 2, stddev 0.01
dropout	dropout rate 0.5
conv (zero padding)	5 x 5 x 32, stride 2, stddev 0.01
max pool	kernel 2 x 2, stride 2
fc	512 x 128, stddev 0.3
fc	128 x 10, stddev 0.3

weights initialization : see the table, as stated in the methodology, the weights should not block the number flowing, so the weights is increasing downstream.

bias initialization : 0.0 for conv, 1.0 for fc

learning rate : 1e-3

batch size : 256


steps : 30k

result :


a. test accuracy : 90.09%

comment : looks fine, the features map works

b. learning curve :



validation accuracy



cross entropy

comment : the learning curve is more reasonable, and the cross entropy is by large dropping below 1.0

c. confusion matrix :

Confusion Matrix of testing data:

```
[[1635  14  12  11  8  4  28  4  9  19]
 [  97 4718  37  61  66 11  15  64 13  17]
 [  16  49 3854  94  28 20  11  39 13  25]
 [  19  55  49 2540  8  68 23 12  37  71]
 [  23  68  16  24 2324 10  15  13  8  22]
 [  13  19  14  101 16 2138 42  5  14  22]
 [  49  16  3  21  22  55 1760  1  37  13]
 [  13  66  41  24  17  6  6 1837  2  7]
 [  27  13  12  45  14  13  75  4 1409  48]
 [  34  7  33  15  7  18  11  4  23 1443]]
```

Classification report of testing data:

	precision	recall	f1-score	support
0	0.85	0.94	0.89	1744
1	0.94	0.93	0.93	5099
2	0.95	0.93	0.94	4149
3	0.87	0.88	0.87	2882
4	0.93	0.92	0.92	2523
5	0.91	0.90	0.90	2384
6	0.89	0.89	0.89	1977
7	0.93	0.91	0.92	2019
8	0.90	0.85	0.87	1660
9	0.86	0.90	0.88	1595
avg / total	0.91	0.91	0.91	26032

comment : position where the value is higher than 90, precision or recall that are lower than 0.9 in marked red

conclusion :

Compared with the previous model, this architecture fixed the '1' vs. '4' issue (number in green). It seems feature extraction works. But the model still struggling with '0' vs. '1', '2' vs. '3' and '5' vs. '3' (in red). Therefore, adding another convolutional layer is still promising.

3. Another Convolutional Layer

architecture :


conv (zero padding)	5 x 5 x 32, stride 2, stddev 0.01
conv (zero padding)	5 x 5 x 64, stride 2, stddev 0.02
dropout	dropout rate 0.5
conv	3 x 3 x 256, stride 2, stddev 0.04
max pool	kernel 2 x 2, stride 2
dropout	dropout rate 0.5
fc	1024 x 128, stddev 0.3
fc	128 x 10, stddev 0.5

weights initialization : stated in the table
 bias initialization : 0.0 for conv, 1.0 for fc
 learning rate : 1e-3
 batch size : 256
 steps : 30k


result:

a. test accuracy : 94.03%

b. learning curves



validation accuracy



cross entropy

c. confusion matrix

Confusion Matrix of testing data:

```
[[1685  12   4   2   5   6  12   3   4  11]
 [  66 4799  24  15  51  10  17  97  14   6]
 [   9   34 3950  23  16   5  11  78   4  19]
 [   7   40  22 2601  11  53  21  41  65]
 [  11  46   7   6 2412   4   6  14   2  15]
 [   8  11   9  45   6 2233  41   6   9  16]
 [  39  11   3   9   4  30 1858   2  15   6]
 [   4  31  11   1   3   3   1 1961   1   3]
 [  17   9  10  17   5  11  74   5 1473  39]
 [  29   6   9   6  12   7   4   6   9 1507]]
```

Classification report of testing data:

	precision	recall	f1-score	support
0	0.90	0.97	0.93	1744
1	0.96	0.94	0.95	5099
2	0.98	0.95	0.96	4149
3	0.95	0.90	0.93	2882
4	0.96	0.96	0.96	2523
5	0.95	0.94	0.94	2384
6	0.91	0.94	0.92	1977
7	0.89	0.97	0.93	2019
8	0.94	0.89	0.91	1660
9	0.89	0.94	0.92	1595
avg / total	0.94	0.94	0.94	26032

comment : red is where the value is higher than 60, or precision or recall is lower than 0.9

conclusion :

Again we can look at the confusion matrix to see how feature extraction works in the recognition. First, all the position marked red in the previous experiment is all fixed, but a new issue on '1' vs. '7' is introduced, which is very natural in real life.

I used CPU to train, so I decided to stop adding more convolutional layer. Next I will tune the weight initialization, and learning rate.

4. Weights Initialization

overview : Weights initialization is important and an active fields in research. This experiment didn't use too many mathematical theory to guide the process. The idea is just tuning the standard deviation. Higher stddev means the weights have more chance to distributed in a larger interval, so that weights are less likely to block the information flow because they are too small. Basically, the stddev drops while going upstream of the network.

weight initialization :

conv 1	0.01
conv 2	0.05
conv 3	0.07
fc 1	0.1
fc 2	0.2

other : same as experiment 3

result :

a. test accuracy : 94.94%

b. learning curve : skipped no more information than experiment 3

c. confusion matrix

Confusion Matrix of testing data:										Classification report of testing data:				
											precision	recall	f1-score	support
[[1704 12 2 4 1 4 7 3 3 4]										0	0.89	0.98	0.93	1744
[67 4888 23 17 25 13 10 34 14 8]										1	0.96	0.96	0.96	5099
[9 27 4027 29 9 4 4 17 11 12]										2	0.97	0.97	0.97	4149
[18 32 18 2633 2 48 18 4 45 64]										3	0.95	0.91	0.93	2882
[14 47 11 9 2405 1 9 10 4 13]										4	0.98	0.95	0.96	2523
[7 10 9 39 4 2251 39 2 13 10]										5	0.95	0.94	0.95	2384
[44 11 2 3 4 28 1859 1 20 5]										6	0.94	0.94	0.94	1977
[6 63 28 8 7 6 0 1895 1 5]										7	0.96	0.94	0.95	2019
[19 13 11 13 5 5 40 1 1533 20]										8	0.93	0.92	0.93	1660
[31 10 11 6 1 7 2 1 6 1520]]										9	0.92	0.95	0.93	1595
										avg / total	0.95	0.95	0.95	26032

conclusion :

The architecture is promising to be my best model. Weights initialization works well here and raised the test accuracy by 0.91%. A relatively good improvement.

5. My Best Model

Based on experiment 4, I introduced exponential learning rate decay and get 95.25% on test set.

learning_rate = $1e-3 * 0.98^{(global_step/500)}$.

The confusion matrix is below :

Confusion Matrix of testing data:										Classification report of testing data:				
											precision	recall	f1-score	support
[[1693 7 1 4 3 2 11 6 4 13]										0	0.90	0.97	0.93	1744
[62 4856 23 18 36 11 7 65 17 4]										1	0.97	0.95	0.96	5099
[6 21 4003 35 19 1 7 32 9 16]										2	0.98	0.96	0.97	4149
[14 27 19 2665 8 34 18 10 46 41]										3	0.95	0.92	0.94	2882
[17 30 12 2 2432 2 4 10 6 8]										4	0.97	0.96	0.96	2523
[7 13 9 44 6 2250 35 2 8 10]										5	0.97	0.94	0.96	2384
[37 13 3 5 5 14 1877 3 15 5]										6	0.94	0.95	0.94	1977
[4 37 14 7 3 2 1 1950 0 1]										7	0.94	0.97	0.95	2019
[25 8 2 14 6 3 35 1 1546 20]										8	0.93	0.93	0.93	1660
[25 7 13 6 2 4 2 3 10 1523]]										9	0.93	0.95	0.94	1595
										avg / total	0.95	0.95	0.95	26032

conclusion :

Stochastic Gradient Descent is a slow method. As we can see in experiment 3, loss is unstably decreasing, and the final loss is not the minimum of the curve. Learning rate decay decreases the footage when the the model approaching the global minimum so that it avoid stepping over the optimal.


IV. Result

Model Evaluation and Validation


a. The confusion matrix is shown in experiment 5.

b. Weights and Biases


Mean/FC_7/Biases




Mean/FC_7/Weights




Mean/FC_8/Biases




Mean/FC_8/Weights




Stddev/FC_7/Biases




Stddev/FC_7/Weights



Stddev/FC_8/Biases




Stddev/FC_8/Weights




(Note: fc7 is fc1, fc8 is fc2, because I assigned indices along the net)

c. convolutional layers : (these may make no sense)



conv 1 before relu



conv 1 after relu

Justification

Compared with the benchmark model, mine lost about 2% in test accuracy. In terms of architecture, I just maxout on the last conv and use stride 2 kernels. In the last conv, I use 3 x 3 instead of 5 x 5 kernel.

V. Conclusion

Reflection

Stories :

It takes me a long time to finish this project. Picking deep learning as my capstone is quite ambitious, especially when you don't have any GPU support. At the beginning, I want to make the app or at least build the multi-digit model. But when it took my laptop 2-3 hours to run a epoch. I think I should just pick the single digit one before I update my hardware.

I started my journey on deep learning by taking the deep learning course on Udacity. However, I quickly found the course is not enough for me to fully understand how convolutional works. My first few models didn't work at all. So I turn to [cs231n](https://cs231n.github.io/) by stanford and apply the techniques and principle to build my model. And it worked!

Learned :

1. Math Is the Most Important.

In TensorFlow, the math is well packaged. In some cases, works for developers are just stacking nodes and layers, which is exactly what I decided to do at the beginning. However, I continuously came across various unexpected issues like adding layers didn't help improve the model. It was math that helps me out when I understood the idea of information flow along the network, which was my golden guide while doing refinements.

2. Documentation is Essential in Development

When we are tool users, documentations help us quickly master the tools. Documentation with examples can minimize the chances we get stuck in development. When we are tool developers, we should always keep in mind that we are not alone in creativity. Tools are valuable only when many people can easily use it.

3. Read More Papers

There are more and more papers in technology. Reading papers help keep track of the latest achievement people make. It's necessary to refresh our mind often today.

Improvements

1. Tensorboard Structure :

The visualization is bad, the summary is not neat.

2. Make an App :

Build an Android App based on the model.